

# Recommending Interesting Writing

Rohan Bansal

The Browser

rohan@thebrowser.com

Uri Bram

The Browser

uri@thebrowser.com

Robert Cottrell

The Browser

robert@thebrowser.com

Jaan Altosaar

Princeton University

altosaar@princeton.edu



Figure 1: Seattle Mariners at Spring Training, 2010.

## ABSTRACT

We describe a system for recommending nonfiction writing to editors at The Browser, a curation service for interesting writing. The editors' goal is to filter new articles from many RSS feeds to choose interesting nonfiction writing to share with readers. To aid the editors, we build a recommender system that classifies articles based on their content. The recommendation model is RANKFROMSETS (RFS), chosen for its scalability and explainability, with architectures that allow editors to understand which words in an article informed a recommendation [1]. Further, editors can choose which latent features of articles to upweight. We show that RFS performs well in classifying articles when evaluated on historical data, and conduct an online evaluation with qualitative feedback from the editors to show that the system performs well in practice on unseen articles. Due to resource constraints, we deploy RFS using a microservices architecture on a cloud computing platform. For reproducibility we open source the end-to-end pipeline and release a demo<sup>1</sup>, training and deployment scripts, and model parameters.<sup>2</sup>

<sup>1</sup><https://the-browser.github.io/recommending-interesting-writing/>

<sup>2</sup><https://github.com/the-browser/recommending-interesting-writing>

## CCS CONCEPTS

• **Applied computing** → Document searching; • **Computing methodologies** → Learning from implicit feedback.

## KEYWORDS

content-based recommendation, open source, user interface

## ACM Reference Format:

Rohan Bansal, Uri Bram, Robert Cottrell, and Jaan Altosaar. 2018. Recommending Interesting Writing. In *KDD '20: International Workshop on Industrial Recommendation Systems, August 23–27, 2020*. ACM, New York, NY, USA, 1 page. <https://doi.org/10.1145/1122445.1122456>

## REFERENCES

- [1] Jaan Altosaar, Wesley Tansey, and Rajesh Ranganath. 2020. RankFromSets: Scalable Set Recommendation with Optimal Recall. *American Statistical Association Symposium on Data Science & Statistics* (2020).

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from [permissions@acm.org](mailto:permissions@acm.org).

KDD 2020, August 23–27, 2020,

© 2018 Association for Computing Machinery.

ACM ISBN 978-1-4503-XXXX-X/18/06...\$15.00

<https://doi.org/10.1145/1122445.1122456>