


# Recommending Interesting Writing using a Controllable, Explanation-Aware Visual Interface

Rohan Bansal	Jordan Olmstead	Uri Bram	Robert Cottrell
The Browser	The Browser	The Browser	The Browser
rohan@thebrowser.com	jordan@thebrowser.com	uri@thebrowser.com	robert@thebrowser.com

Jaana Altsaara  
Princeton University  
altsaara@princeton.edu



**Figure 1:** End-to-end pipeline for recommending writing to editors at The Browser. The RANKFROMSETS recommendation model [2] is trained on data consisting of positive examples and negative examples collected from the editors’ history of curated articles, and negative examples from news sources. After training and offline evaluation of the recommendation model, RANKFROMSETS is deployed as a microservice on Amazon Web Services Lambda, with the visual interface hosted on Github Pages. Editors can control the recommender system using the visual interface, which can aid in their decision-making. The editors’ interrogation of the recommendation model and feedback is then used to inform of data collection and training.

## ABSTRACT

We build a visual interface for recommending articles to editors at The Browser, a curation service for interesting writing. From a large list of candidates, editors decide which articles are selected and shared with subscribers. To aid the editors in this decision-making task, we build a visual interface for a recommendation model, RANKFROMSETS (RFS) [2], that classifies articles based on their words. Control of the recommendation model is built into the visual interface. For example, an editor can use a topic slider to receive a new list of recommendations according to topical words in articles. These topic sliders might be used to increase or decrease the ranking of articles with words related to crime, business, or technology. The visual interface is also designed to be explanation-aware: words that contribute positively or negatively to an article’s ranking are displayed. For the backend of the visual interface, RFS is trained on historical data. In an offline empirical study, we find that RFS outperforms BERT [4], a competitive classification model, in terms of recall. Further, we measure RFS to be 10 times faster to

train and to return predictions 2000 times faster than BERT. These are beneficial properties for the visual interface, and we demonstrate that RFS can be deployed on the free tier of AWS Lambda using a short python script and numpy dependency. For reproducibility, transparency, and trust of the visual interface, we open source and release a public demonstration,<sup>1</sup> data collection, training and deployment scripts, and model parameters.<sup>2</sup>

## CCS CONCEPTS

• Applied computing → Document searching; • Computing methodologies → Learning from implicit feedback.

## KEYWORDS

content-based recommendation, open source, visual interface

## ACM Reference Format:

Rohan Bansal, Jordan Olmstead, Uri Bram, Robert Cottrell, and Jaana Altsaara. 2020. Recommending Interesting Writing using a Controllable, Explanation-Aware Visual Interface. In *RecSys ’20: Workshop on Interfaces and Human Decision Making for Recommender Systems, September 22–26, 2020*. ACM, New York, NY, USA, 4 pages. <https://doi.org/10.1145/1122445.1122456>

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from [permissions@acm.org](mailto:permissions@acm.org).

RecSys 2020, September 22–26, 2020,

© 2020 Association for Computing Machinery.

ACM ISBN 978-1-4503-XXXX-X/18/06...\$15.00

<https://doi.org/10.1145/1122445.1122456>

<sup>1</sup><https://the-browser.github.io/recommending-interesting-writing/>

<sup>2</sup><https://github.com/the-browser/recommending-interesting-writing>

## 1 INTRODUCTION

Creative nonfiction, longform journalism, and blog posts are examples of the types of articles curated by The Browser’s team of editors. The editors read a large number of articles from various publications to select content to recommend to subscribers.

In building a recommender system to help editors sift through many documents, it is motivating to highlight the trade-off in user privacy intrinsic to recommender systems. A machine learning model must exploit information about a user. However, the incentive structures of operating a recommender system within a business can influence decisions around privacy and transparency [5]. For example, business models that rely on online advertising may engender recommender systems that upweight attention-grabbing content and hence time spent looking at ads. Such content might maximize a user’s time spent with a service over time at the expense of long-term user experience or consent. Outside of recommender systems, privacy-preserving and open source tools such as the Signal encrypted messaging service<sup>3</sup> may provide improved user experience in terms of privacy-preserving, transparent, and explainable algorithms and visual interfaces [3]. Due to poor incentive structures for releasing recommender systems and visual interfaces that exploit private information about users, there are few examples of end-to-end, open source, free-to-deploy pipelines for recommending content to users using a visual interface. This serves as motivation to build and deploy a recommendation model and corresponding explanation-aware visual interface to give users control, and inform them about how data is being used to make recommendations.

We build an end-to-end recommender system visual interface to address two aims: (1) to aid editors at The Browser in their decision-making task, and give them control through an explanation-aware interface, and (2) to release a lightweight, performant, open-source visual interface framework for developing and deploying explanation-aware recommender systems for document recommendation. In an offline evaluation, we show that the recommendation model we use for the visual interface outperforms BERT, a competitive document classification model. In a qualitative study, the control and explanations provided by the visual interface help editors in their decision-making and can help find bugs in the recommendation model.

## 2 RECOMMENDATION MODEL

RANKFROMSETS (RFS) is the recommendation model that powers the visual interface; the main part of the pipeline illustrated in Figure 1. RFS scales to large numbers of articles, and can maximize the evaluation metric of recall [1, 2]. Recall, or the fraction of true positives returned by a recommendation model, is an appropriate evaluation metric for recommending interesting writing to editors at The Browser. A recommendation model such as RFS can be readily backtested with recall as an evaluation metric, as historical data contains positive examples (articles selected by the editors) but rarely contains negative examples (articles seen but not selected by the editors). Further, as our goal is to build an explanation-aware visual interface that can also serve to control recommendations,

RFS is fast, interpretable, and simple to integrate into a user interface.

RFS is a recommendation model defined by a binary classifier. For a user  $u$  and item  $m$  with attributes  $x_m$  (the set of unique words in an article), RFS is described by the probability of  $y_{um} = 1$  (user  $u$  consuming item  $m$ ):

$$p(y_{um} = 1 \mid u, m) = \sigma(f(u, x_m)),$$

where  $\sigma$  is the sigmoid function. To parameterize the binary classifier in RFS, we use an inner product architecture:

$$f(u, x_m) = \theta_u^\top \left( \frac{1}{|x_m|} \sum_{j \in x_m} \beta_j \right). \quad (1)$$

In this architecture, the user embedding  $\theta_u$  includes a dimension that is fixed to unity. Word embeddings  $\beta_j$  (including a bias dimension for every word) and the publication embedding are fit with maximum likelihood estimation, and negative examples are sampled uniformly at random to balance positive examples.

## 3 VISUAL INTERFACE

The visual interface is designed using RFS as the backend recommendation model. Using the inner product architecture for RFS makes it easy to build a visual interface that is interpretable to provide explanations for why an item was recommended, and enables control so users can filter recommendations to help with decision-making.

*Explanation-aware recommendation.* The user embedding  $\theta_u$  and word embeddings  $\beta_j$  in Equation (1) can be used to interpret a recommendation. The logit for a given document with a set of words  $x_m$  is the sum of per-word logits, which are computed as the inner product of the user embedding and word embedding. The per-word contribution of a word in a document to the logit that determines the document’s ranking in a list of recommendations is therefore

$$w_{uj} = \theta_u^\top \beta_j. \quad (2)$$

This weight  $w_{uj}$  helps explain why a document was recommended, using information about both the user  $u$  and the word  $j$ . In the visual interface, words in a document are first sorted by their contributions to a document’s logit  $w_{uj}$ , and the top words are displayed. Similarly, words that lower a document’s ranking are also displayed, to inform a user of which words detract from the recommendation of a document.

*Interface for controlling recommendations.* In a decision-making task, a user such as an editor for The Browser may wish to filter recommendations according to topics such as crime, technology, or business. The recommendations output by RFS can be controlled, by altering the per-word contributions in Equation (2) according to whether a word is topical. This is accomplished by first calculating words related to a topic word using pre-trained word embeddings from BERT [4, 7]. Words related to a topic are defined by a heuristic: the cosine similarity between all words and a topic word such as ‘business’ are computed, and the top 15 words closest in cosine distance are stored as topical words. Then, a slider in a visual interface

<sup>3</sup><https://signal.org/>

is used to increase or decrease the per-word contributions of topical words to a document’s logit. Let the user-input slider value be  $\alpha$ , and the set of topical word indices be  $T$ . Then the user-controlled version of Equation (1) becomes

$$f(u, x_m) = \theta_u^\top \left( \frac{1}{|x_m|} \sum_{j \in x_m} (1 - \mathbb{I}[j \in T])\beta_j + \mathbb{I}[j \in T]\alpha \text{sgn}(w_{uj})\beta_j \right). \quad (3)$$

The sign function  $\text{sgn}(\cdot)$  is applied to the per-word contribution to a document’s logit. This is included since a word might contribute negatively to a document’s logit, yet a user may wish to increase the weight of a related topical word.

## 4 EVALUATION

We conduct an offline empirical study of the performance of RANKFROMSETS to assess its performance as a recommendation model. Then we qualitatively evaluate the visual interface to see whether the explanation-aware, controllable interface enabled by RFS can help make editors at The Browser make better decisions.


*Data Collection and Preprocessing.* For positive examples, we use the historical set of articles curated by editors at The Browser. We augment the training data with articles selected by the editors of other curation services, and treat all positively-labeled examples curated by editors as data from a single user due to a paucity of data. We use articles from news websites as examples with negative labels, and collect additional articles with negative labels from websites most-featured by the editors to mimic the editorial process of reading a large swath of articles in a feed and distilling an article list to a select few. For preprocessing the data we use the tokenizer released by Devlin et al. [4] and discard words not recognized by the tokenizer. This procedure results in a dictionary with 30k words, and 646k datapoints with 27k positive labels.

*Metrics.* Performance of the recommendation models is assessed with recall, and 15% of the datapoints are held out for the validation and test sets respectively.

*Experimental setup: RankFromSets.* We cross-validate using the RMSProp optimizer [6] with a momentum of 0.9 and grid search over learning rates of  $\{10^{-2}, 10^{-3}, 10^{-4}, 10^{-5}\}$ , whether or not to initialize from pre-trained BERT embeddings [7], and embedding sizes of  $\{10, 25, 50, 100, 500, 1000\}$ . This model is trained on an NVIDIA Tesla P100 GPU.

*Experimental setup: BERT.* To fine-tune BERT, we use the AdamW optimizer with a linear learning rate scheduler and warmup steps, with a batch size of 32 and maximum input length of 512 as in Devlin et al. [4] and Wolf et al. [7]. A grid search is performed over learning rates of  $\{2, 3, 4, 5\} \times 10^{-5}$ , warmup steps of  $\{10^2, 10^3, 10^4\}$ , and total training steps of  $\{10^2, 10^3, 10^4, 10^5\} \times 5$ . The model is trained on an NVIDIA Tesla V100 GPU.

The best-performing model of RFS is selected for deployment, and recall is evaluated on the test set, after using early stopping according to validation recall. The results are shown in Table 1, and RFS outperforms BERT by 14%. Furthermore, RFS achieves better performance ten times faster than BERT, as shown in Figure 2. In a test to measure the speed of recommending  $10^4$  held-out articles, RFS



**Figure 2: RANKFROMSETS achieves better performance faster than BERT in terms of validation recall during training.**

Recommendation Model	Recall @ 1000 (%)
RANKFROMSETS	53.1
BERT	46.6

**Table 1: RFS outperforms BERT in an offline evaluation, on a task of predicting which articles editors at The Browser would feature based on words in the articles.**

ranked all articles in 120 ms on a CPU, while BERT took 4 m 54 s to rank the articles on an NVIDIA Tesla V100 GPU. This represents a 2000-fold improvement in speed, which is beneficial for the controllable visual interface that requires Equation (3) to be quickly computed in response to user input.

*Qualitative Evaluation.* In a user study, editors at The Browser provided feedback that they used the visual interface to choose articles, and found this to be an improved workflow. The control over recommendations, and explanation-aware visual interface provided by RFS helped elicit bugs in data collection (such as foreign language sources, or fiction writing) and provides an enjoyable user experience.

## 5 DEPLOYMENT

We deploy the visual interface using Github Pages, and deploy the backend recommendation model, RFS, as a microservice hosted on Amazon Web Services Lambda. As Equation (3) is cheap to compute, the lambda function is a short python script that simply requires numpy as a dependency, compared to BERT which would require a hosted GPU solution.

## 6 DISCUSSION

We built a visual interface for a recommender system powered by RFS, a flexible recommendation model. Empirically, we demonstrated that RFS outperforms BERT in an offline evaluation, while being orders of magnitude faster during training and recommendation. By deploying RFS to AWS Lambda and hosting the visual interface on Github Pages, we demonstrated a fully open-source pipeline for creating an explanation-aware, controllable visual interface for document recommendation for editorial decision-making. Future work includes studying whether the transparency and control provided by open-source recommendation systems can improve user experience and educate users in how their information is exploited remains an open problem.

## ACKNOWLEDGMENTS

The authors are grateful to Christian Bjartli for help with data collection.

## REFERENCES

- [1] Jaan Altosaar. 2020. *Probabilistic Modeling of Structure in Science: Statistical Physics to Recommender Systems*. Ph.D. Dissertation. Princeton University.
- [2] Jaan Altosaar, Wesley Tansey, and Rajesh Ranganath. 2020. RankFromSets: Scalable Set Recommendation with Optimal Recall. *American Statistical Association Symposium on Data Science & Statistics* (2020).
- [3] K. Cohn-Gordon, C. Cremers, B. Dowling, L. Garratt, and D. Stebila. 2017. A Formal Security Analysis of the Signal Messaging Protocol. In *2017 IEEE European Symposium on Security and Privacy (EuroSP)*. 451–466.
- [4] Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. 2019. BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding. In *Association for Computational Linguistics*. Association for Computational Linguistics, Minneapolis, Minnesota, 4171–4186. <https://doi.org/10.18653/v1/N19-1423>
- [5] Nicholas Diakopoulos. 2020. *Oxford Handbook of Ethics and AI*. Oxford University Press, Chapter Accountability, Transparency, and Algorithms.
- [6] Tijmen Tieleman and Geoffrey Hinton. 2012. Lecture 6.5-rmsprop: Divide the gradient by a running average of its recent magnitude. *COURSERA: Neural Networks for Machine Learning* (2012).
- [7] Thomas Wolf, Lysandre Debut, Victor Sanh, Julien Chaumond, Clement Delangue, Anthony Moi, Pierric Cistac, Tim Rault, R’emi Louf, Morgan Funtowicz, and Jamie Brew. 2019. HuggingFace’s Transformers: State-of-the-art Natural Language Processing. *ArXiv abs/1910.03771* (2019).